

SLR Project P4 — Distributed MapReduce Final Documentation

Project: Custom distributed MapReduce in Go over remote lab machines (`*.enst.fr`)

Repository: `slr_map_reduce`

Last updated: 2026-07-01

1. Executive Summary

This project implements a fault-tolerant batch MapReduce framework with:

- A Go orchestrator (`scripts/mapreduce`) that deploys workers, drives phases, and merges outputs.
- HTTP workers (`scripts/worker`) that execute map/reduce, partition intermediate data by hash, and exchange buckets.
- Pluggable jobs implementing `mrjob.Mapper` / `mrjob.Reducer` (`scripts/jobs/*`).
- A Kafka Streams baseline (`kafka/*`) to compare compute time and operational trade-offs.

Three additional use-case jobs were implemented and validated end-to-end:

1. `langdetect` — language distribution via stop-word scoring.
2. `domainpop` — domain popularity from `WARC-Target-URI` headers.
3. `docdensity` — document size and character-density counters.

All reducers emit integer strings to remain compatible with the orchestrator's integer-sum merge.

2. Architecture

2.1 Components

Component	Path	Responsibility
Orchestrator	<code>scripts/mapreduce</code>	Build worker with selected job, deploy to hosts, run load→map→reduce→collect, merge final output
Worker	<code>scripts/worker</code>	Expose HTTP API, map local chunks/files, write hash-partitioned buckets, fetch peers, reduce, return result
Job API	<code>scripts/mrjob/api.go</code>	Defines <code>KeyValue</code> , <code>Mapper</code> , <code>Reducer</code> interfaces
Jobs	<code>scripts/jobs/*</code>	Independent map/reduce logic selected with <code>-job</code>
Host discovery	<code>scripts/make_hosts</code>	Build runtime host lists from lab availability
Kafka baseline	<code>kafka/*</code>	Rootless KRaft deploy + cleanup + Streams wordcount driver

2.2 End-to-end lifecycle

1. Parse flags (`-hosts`, `-job`, `-input` or `-commoncrawl`, `-n`, etc.).
2. Build worker binary with generated job binding (`job_binding_generated.go`).
3. Deploy the worker binary only to the initial active hosts and start remote `mr_worker` with PID tracking.
4. Health-check those active workers (`/health`) until `N` slots are filled; keep the remaining discovered hosts as cold spares.
5. Load inputs (`POST /data` for local chunks, `POST /load` for WET URLs).
6. Run map (`POST /map`) where each worker writes bucket files partitioned by reducer.
7. Run reduce (`POST /reduce`) where each reducer pulls exactly one bucket from each peer.
8. Collect (`GET /result`) and merge counts in orchestrator.
9. If a slot host fails, deploy and health-check a cold spare on demand before replaying the needed phases for that slot.
10. Cleanup remote workers on every host that was actually started.

2.3 Why integer outputs are mandatory

The orchestrator merge path sums values via `strconv.Atoi` (`coordinator.mergeResults`). Therefore job reducers must emit integer strings.

- `wordcount`, `langdetect`, `domainpop`: emit count occurrences.
- `doccdensity`: emits integer counters (`stat:*`, `hist:*`) so averages/ratios are derived post-run.

This preserves compatibility without changing the merge protocol.

3. Protocol and Dataflow

3.1 Worker HTTP API

Endpoint	Method	Description
<code>/health</code>	GET	Readiness/liveness probe
<code>/data</code>	POST	Upload local input chunk
<code>/load</code>	POST	Download Common Crawl URLs assigned to this worker
<code>/map</code>	POST	Execute map with worker slot ID and peer list
<code>/intermediate?reducer=i&n=N</code>	GET	Return pre-partitioned KV bucket for reducer <code>i</code>
<code>/reduce</code>	POST	Fetch buckets from all peers and run reduce
<code>/result</code>	GET	Return final <code>key\tvalue</code> lines
<code>/debug/pprof/*</code>	GET	Profiling routes when worker starts with <code>-pprof</code>

3.2 Space-Time Diagram (Protocol Execution)

3.3 Hash partitioning

Reducer assignment is deterministic:




Figure 1: Sequence Diagram

$$extreducer(key) = \text{FNV32a}(key) \bmod N$$

N stays fixed for the whole job. Fault-tolerance replaces hosts behind slots, not slot count.

3.4 $O(1)$ shuffle read path

Workers write `map-bucket-<N>-<idx>.jsonl` files during map.

At reduce time, peer fetch is direct file lookup per reducer index:

- no scan over full intermediate set,
- no filter pass,
- one read target per peer request.

This gives $O(1)$ bucket selection per request and $O(N^2)$ total peer requests for full shuffle.

4. Fault Tolerance

4.1 Slot model

- A **slot** is a logical worker index $0..N-1$.
- A **spare** is a discovered host that is not assigned a slot initially; it is deployed on demand only if a slot host must be replaced.
- When spares run out, a healthy worker may temporarily own more than one slot so N stays fixed.
- Slot state machine: **pending** $\rightarrow$ **loaded** $\rightarrow$ **mapped** $\rightarrow$ **reduced** $\rightarrow$ **done**.
- Replacing a failed slot host resets that slot to replay required prerequisites.

4.2 Failure handling

1. Background health watcher probes `/health` on active slot hosts and marks hosts dead after repeated failures.
2. Phase calls retry with exponential backoff and bounded attempts.
3. Reduce failures that indicate `fetch from slot X peer ...` are attributed to slot **X**, even if multiple slots currently share the same physical host.
4. When a slot host fails, the coordinator first tries to deploy a cold spare on demand; if none can be activated, it falls back to rebinding the slot onto a surviving healthy worker.
5. On slot-host replacement, coordinator epoch increments and previously reduced slots are demoted to mapped (reduce replay).

This prevents stale peer topology from leaking into final reduce outputs.

Atomic writes status:

- WET file downloads: `CreateTemp` + `os.Rename` — implemented ✓
- Worker binary deploy: `mv` on same filesystem — implemented ✓
- Final output: written directly to the output file path (no temp+rename). If any worker fails during the final result collection phase, the orchestrator exits with no output at all. [known gap]

5. Implemented Jobs

Because Common Crawl is one of the primary training data sources for large language models (GPT, LLaMA, Mistral, and many others), these jobs represent the kind of analysis that precedes large-scale dataset curation — identifying language composition, domain provenance, and content quality before deciding what to include in a training corpus.

5.1 langdetect (`scripts/jobs/langdetect`)

Map logic

- Skip blank/WARC metadata lines.
- Tokenize letters, lowercase.
- Score against stop-word sets (`en`, `fr`, `de`, `es`, `it`).
- Emit one key `lang:<code>` with value 1 per classified content line.

Reduce logic

- Count values per language key.

Output example

- `lang:en\t12345`

5.2 domainpop (`scripts/jobs/domainpop`)

Map logic

- Primary: parse `WARC-Target-URI: <url>` lines.
- Fallback: extract absolute `http://` / `https://` URLs in plain text.
- Normalize host (lowercase, strip port and leading `www.`).
- Emit `<domain>\t1`.

Reduce logic

- Count values per domain.

5.3 docdensity (`scripts/jobs/docdensity`)

Map logic (per non-empty content line)

- `stat:lines -> 1`
- `stat:chars.total -> <line rune count>`
- `stat:chars.alnum -> <letters+digits>`
- `stat:chars.other -> <non-alnum>`
- `hist:len.<bucket> -> 1`

Reduce logic

- Sum integer values per key (not only count occurrences).

Derived metrics

$$extavglinelength = \frac{\text{stat:chars.total}}{\text{stat:lines}}, \quad extalnumdensity = \frac{\text{stat:chars.alnum}}{\text{stat:chars.total}}$$

6. Job Execution Results (Executed on 2026-06-14)

6.1 Execution environment and method

In this machine/session, direct remote deployment over lab SSH is not guaranteed. To still execute the real runtime path, each job was run through the actual worker HTTP pipeline locally:

1. Generate a temporary worker build with injected job binding.
2. Start worker on local port.
3. POST `/data` with input.
4. POST `/map` with peers list containing only itself.
5. POST `/reduce`.
6. GET `/result` and persist output under `run/local_jobs`.

This exercises real map/reduce/job code and protocol handlers, excluding only SSH/SCP deployment.

6.2 Inputs used

- `langdetect, docdensity: test_input.txt` (15 English lines).
- `domainpop`: synthetic WARC-like lines containing multiple `WARC-Target-URI` entries.

6.3 Captured outputs

A) langdetect Output file: `run/local_jobs/langdetect_result.txt`

```
lang:en 15
```

Interpretation: all 15 lines were classified as English, matching the input corpus.

B) domainpop Output file: `run/local_jobs/domainpop_result.txt`

```
example.com 2
golang.org  1
news.ycombinator.com 2
```

Interpretation: domains and counts match inserted `WARC-Target-URI` records.

C) docdensity Output file: run/local_jobs/docdensity_result.txt

```
hist:len.00000-00079    15
stat:chars.alnum      386
stat:chars.other       73
stat:chars.total      459
stat:lines    15
```

Derived metrics (computed with LC_ALL=C awk):

- avg_line_length = 30.600
- alnum_density = 0.840959

Interpretation: all lines fall in <80 bucket; dataset is predominantly alphanumeric text.

7. Performance and Bottleneck Audit

7.1 Amdahl's Law and Performance Metrics

Experiment Setup: Fixed workload of 8 WET files, measuring Map+Reduce+Collect. The same dataset was used for all point measurements to properly validate Amdahl's Law. The wordcount job on this corpus produced approximately **24 million unique words**.

- **1 node:** 206.95s (Baseline)
- **16 nodes:** 40.80s (Speedup 5.07x)
- **32 nodes:** 37.94s (Speedup 5.45x — peak)
- **64 nodes:** 39.06s (Speedup 5.30x — regression)

The speedup rises from 1 to 32 nodes, peaks at $S \approx 5.45$, then regresses slightly at 64 nodes ($S \approx 5.30$). This illustrates all three phases of Amdahl's law: strong gains from 1–8 nodes (parallel work dominates), diminishing returns from 8–32 nodes (communication overhead grows), and regression at 64 nodes (synchronization overhead exceeds parallelization gains).

Serial Fraction Calculation: Using the peak value $S(32) \approx 5.45$, solving $1/f = 5.45$ gives a serial fraction $f \approx 18\%$. This serial portion stems from the orchestrator's collect phase (sequential merge of N results) and the $N \times (N - 1)$ HTTP connections during shuffle at 64 nodes.

7.2 Confirmed code-path bottlenecks

1. Map-phase memory amplification

- `runMap` builds `map[string][]string` with one string per emitted value.
- `handleMap` then copies into per-reducer `[] []KeyValue` buckets before writing.
- Effect: high temporary allocations and duplicated resident intermediate data.

2. Shuffle serialization overhead

- `/intermediate` reads full JSONL bucket to slice and re-encodes JSON array.
- Reducer fetch decodes full array in memory.
- Effect: extra encode/decode CPU and extra memory pressure.




Figure 2: Amdahl's Law Graph

3. Aggressive scanner buffer size

- Scanner buffers are configured to 64MB globally in map/intermediate readers.
- Effect: high baseline memory per concurrent reader.

4. Serial final merge

- Coordinator merge and global sort are single-threaded.
- Effect: contributes to Amdahl serial fraction at larger N.

7.3 Existing mitigations already in code

- Node-local temp directories (`os.MkdirTemp`) avoid NFS for intermediate and downloaded input files.
- Pull-based bucket fetch keeps remote read selection $O(1)$.
- Slot replacement + epoch logic protects consistency under failures.

7.4 Profiling instrumentation

Worker profiling is available with `-pprof` (added in this delivery).

Orchestrator pass-through: `-worker-pprof`.

CPU profile (30s)

```
go tool pprof "http://<worker>:9090/debug/pprof/profile?seconds=30"
```

Heap profile

```
go tool pprof "http://<worker>:9090/debug/pprof/heap"
```

Recommended optimization order:

1. Stream intermediate bucket responses (`io.Copy JSONL`) instead of array re-encode.
2. Introduce combiner/in-mapper aggregation for associative jobs.
3. Lower scanner max size and make it configurable.
4. Parallelize or tree-merge final aggregation in coordinator.

8. Kafka Streams and Hadoop Comparison Documentation

8.1 Three-Way Comparison

Feature	Our MapReduce	Hadoop	Kafka Streams
Model	Batch, 1 stage	Batch, multi-stage	Stream, unbounded
Storage	/tmp, no replication	HDFS 3x replication	Replicated broker log
Data locality	No — HTTP pull	Yes — run where data lives	Broker partitions
Stragglers	Not handled	Speculative execution	N/A
8-node wordcount	0.27s	not run	10.1s

8.2 Batch processing vs. stream processing

Our system and Hadoop are *batch* frameworks: they process a bounded, fixed input, produce a final result, and stop. Kafka Streams is a *streaming* framework: it processes an unbounded sequence of records continuously, updating output state stores in real time. These are fundamentally different execution models suited to different workloads.

	Batch (MapReduce)	Stream (Kafka Streams)
Input	Bounded, known upfront	Unbounded, arrives continuously
Output	Written once at end	Continuously updated state
Latency	Minutes to hours	Milliseconds to seconds
Fault tolerance	Re-run affected tasks	Replicated log replay
Use case	Historical analysis	Real-time monitoring

8.3 Our system vs. Hadoop

Hadoop is a production-grade Java implementation of the same Dean & Ghemawat MapReduce model. The table compares the two on dimensions relevant to this project.

Aspect	Our system	Hadoop (HDFS + YARN)
Storage	Intermediate files on /tmp; no replication	HDFS: blocks replicated 3× across DataNodes
Input distribution	HTTP pull from Amazon S3 or orchestrator push	Files stored on HDFS; map tasks run <i>where the data is</i> (data locality)
Fault tolerance	Re-run slot on spare; epoch-invalidated reduces	Task re-attempt (up to 4×); speculative execution for stragglers
Scheduling	Fixed N slots; round-robin URL assignment	YARN ResourceManager: dynamic container allocation, multiple queues, pre-emption
Combiner	Not implemented	Optional local combiner reduces shuffle volume
Second-stage reduce	Not implemented	Supported: reducer output can feed another MapReduce job
Serialisation	TSV (text)	Writable (binary), SequenceFile, Avro, ORC, Parquet
Language	Go; job compiled into worker binary	Java; JAR submitted at runtime
Lines of code	≈ 3,000	≈ 10,000,000
No-root install	Yes	Requires cluster setup (HDFS, YARN daemons)

Major Differences vs. Hadoop:

- **Data Locality:** Hadoop uses HDFS to run tasks where data is stored, eliminating network transfer at map time. Our system pulls all input at load time via HTTP.
- **Speculative Execution:** Hadoop detects stragglers and launches backup task attempts. In our system, one slow node blocks the entire phase.

8.4 Our system vs. Kafka Streams (measured)

In a direct comparison using the 8-file Common Crawl corpus on 8 nodes, the Go MapReduce system completed in **0.266 s**, whereas Kafka Streams took **10.095 s**. The difference is architectural: Kafka persists every record to a replicated broker log (durability the Go system does not offer), paying an I/O cost per record. On small bounded inputs, this persistence overhead dominates.

In a separate 50-node benchmark, Kafka Streams completed reliably in **248.5 s**, whereas the Go system struggled with SSH connection limits at that scale. Kafka manages partitions dynamically and is inherently more robust for large-scale or continuous workloads.

System	8-node compute (s)	50-node compute (s)	Durability
Go MapReduce	0.266	unstable	None
Kafka Streams	10.095	248.5	Replicated log

Timing methodology: the Go orchestrator reports `TIMING ... compute_seconds=...` for `map+reduce+collect`; the Kafka runner reports `TIMING compute_seconds=...` when consumer lag reaches zero. In both paths, input staging happens before the timer starts, making the two values directly comparable.

8.5 Delivered assets

- `kafka/deploy_kafka.sh`: rootless KRaft deployment (no Docker, no root).
- `kafka/clean_kafka.sh`: full cleanup and process termination.
- `kafka/WordCountJob.java`: Streams wordcount with self-timed completion.
- `kafka/run_kafka_wordcount.sh`: end-to-end benchmark driver.

8.6 Fair benchmark protocol

1. Same input file for both systems.
2. Same node count (N).
3. At least 3 runs per point, use median.
4. Report caveat: Kafka includes log persistence semantics absent in volatile MapReduce run.

9. Troubleshooting and Operations

9.1 Pain Points: What went wrong and why

- **NFS I/O Saturation (Biggest Pain Point)**: When all workers tried to read from `/cal/commoncrawl` simultaneously, the NFS became heavily saturated and unresponsive.
 - *Fix*: We transitioned to streaming data directly from `data.commoncrawl.org` (Amazon S3), which bypasses the NFS and scales linearly with node count.
- **Zombie Worker Processes**: Killing the orchestrator (e.g. via `SIGKILL` or a lost SSH session) left orphan worker processes running on remote lab machines, holding their assigned ports. The next run would then fail to start fresh workers on those ports or find stale state.
 - *Fix*: Cleanup was made **idempotent** — workers are identified by a PID file at `/tmp/mr-worker.pid`; the cleanup step kills via PID and removes all `/tmp/mr-worker*` files, tolerating repeated invocations and missing files gracefully.
- **Memory Pressure**: Processing large datasets caused OOM or disk-space issues on smaller clusters because intermediate data is heavily buffered in memory before being written to `/tmp`.

9.2 Common issues

- **scp permission denied on /tmp/mr-worker**
 - Cause: stale file ownership from another user/process.
 - Current mitigation: timestamped deploy temp path + move on remote.
- **No remote workers become healthy**
 - Validate SSH, port reachability, and remote /tmp/mr-worker.log. The wrapper fetches a pool of $2 * N$ hosts, but only the first N successful startups become active workers.
- **Unexpected zero output**
 - Confirm job-specific input expectations:
 - * domainpop expects WARC-Target-URI lines or explicit URLs.
 - * langdetect ignores very short/noisy lines and metadata headers.

9.3 Verification checklist

- `go test ./...` in `scripts/` passes.
 - `curl /health` is ok for each worker.
 - Map/reduce endpoints return 200.
 - Result file contains expected key family (`lang:*`, `domains`, or `stat:*`).
-

10. Reproducibility Commands

10.1 Standard remote execution (lab hosts)

```
# Word count
./mapreduce.sh -job scripts/jobs/wordcount -input test_input.txt -n 10 -output result_wordcount.txt

# New jobs
./mapreduce.sh -job scripts/jobs/langdetect -input test_input.txt -n 10 -output result_langdetect.txt
./mapreduce.sh -job scripts/jobs/domainpop -input test_input.txt -n 10 -output result_domainpop.txt
./mapreduce.sh -job scripts/jobs/docdensity -input test_input.txt -n 10 -output result_docdensity.txt

# docdensity derived metrics
LC_ALL=C awk -F'\t' '
$1=="stat:chars.total"{t=$2}
$1=="stat:chars.alnum">{a=$2}
$1=="stat:lines"{l=$2}
END{
    printf "avg_line_length=%.3f\n", t/l
    printf "alnum_density=%.6f\n", a/t
}' result_docdensity.txt
```

`mapreduce.sh` fetches $2 * N$ hosts into `hosts.txt` so the orchestrator can keep cold spares available while still running with only N active workers.

10.2 Local single-worker protocol validation (used in this report)

```
# Example for langdetect
cd scripts
BUILD=$(mktemp -d .mr-local-langdetect-XXXX)
cp worker/*.go "$BUILD"/
rm -f "$BUILD"/*_test.go "$BUILD"/builtin_job.go
cat > "$BUILD"/job_binding_generated.go <<'EOF'
package main

import (
    "scripts/mrjob"
    job "scripts/jobs/langdetect"
)

func loadInjectedJob() (mrjob.Mapper, mrjob.Reducer, error) {
    return job.NewMapper(), job.NewReducer(), nil
}
EOF
go build -o /tmp/mr-worker-langdetect "$BUILD"
cd ..

/tmp/mr-worker-langdetect -port 9911 &
curl -X POST --data-binary @test_input.txt http://127.0.0.1:9911/data
curl -X POST -H 'Content-Type: application/json' \
    -d '{"id":0,"peers":["127.0.0.1:9911"]}' http://127.0.0.1:9911/map
curl -X POST http://127.0.0.1:9911/reduce
curl http://127.0.0.1:9911/result
```

11. Conclusion

The project now includes a robust distributed MapReduce framework, three additional analytical jobs, profiling hooks, and Kafka comparison tooling with reproducible scripts.

The three new jobs were executed successfully via the real worker protocol in this environment, and their outputs match expected semantics:

- **langdetect**: full-English corpus detection,
- **domainpop**: accurate domain counts from WARC metadata,
- **docdensity**: consistent aggregate counters and derived ratios.

Next high-impact engineering steps are memory/shuffle streaming optimizations and benchmark completion on the full remote lab pool for final Amdahl and Kafka comparison plots.

The remaining known gaps are: speculative backup tasks, a second MapReduce stage for multi-pass jobs and large-scale empirical evidence beyond 64 nodes.